

Rubric — Track 2: Advanced (Automation Rules Engine)

Total: 100 points; bonus up to +15.

Category	Weight	Excellent	Adequate	Missing
CLAUDE.md quality	10	Project CLAUDE.md captures rule tree shape, pure vs. side-effect boundaries, and worker run steps; reflection is substantive and cites concrete divergences	Present but generic; reflection reads as summary rather than critique	Missing, or unchanged /init output submitted verbatim
Design brief	5	Nested condition builder wireframe is convincing (nested groups visible, add-child affordances shown); two error states designed with copy	Present but incomplete; one wireframe or one error state missing	Missing
Plan + plan-review	10	Plan is dense, ordered by dependency, and lists risk per step; review critiques concretely (line-level suggestions, not "looks good"); student integrated critique or documented why they did not	Plan present, review shallow (< 5 concrete suggestions)	Review missing, or plan is a bulleted list without dependencies
Architecture + arch-review	15	Data model, event flow, worker choice justified against an alternative, retry policy specified numerically; review checks plan consistency step-by-step	Missing one of these components	Missing two or more, or architecture contradicts plan without note

Category	Weight	Excellent	Adequate	Missing
Feature correctness	20	Matcher handles nested trees with depth guard; actions run in order and one failing action does not block others; webhook retries with exponential backoff on a worker; dry-run works and does not persist	Missing dry-run, or missing retry, or matcher handles only flat trees	Core matcher broken, or actions run in wrong order
TDD evidence + coverage	15	Failing-test commits present on all three targets; coverage report shows 85%+ on matcher and executor; retry test asserts both count and intervals	Some TDD, low or missing coverage report; retry test asserts count only	No TDD evidence in git log, or no coverage report
Subagent orchestration	10	At least two distinct subagents used; explicit review loop where one subagent reviewed another's output; artifacts committed	One subagent, no review loop	No subagent evidence in report or git log
Plugin usage	5	superpowers or cavern invoked and visibly shaped the architecture doc; evidence in the workflow report (screenshot or excerpt)	Plugin listed but influence on architecture unclear	No plugin used
HTML workflow report	10	All 8 sections present, screenshots included, coverage report linked	Sections thin or one missing	Report missing

Bonus (+15 max)

- Rule execution history log at `/rules/{id}/history`: +5
- Admin webhook delivery view at `/admin/webhooks`: +5
- Dry-run supports hand-typed event JSON: +5

Grading Notes for Instructor

- The subagent orchestration category is the differentiator between this track and Track 1. If evidence is absent, cap this category at 3 points regardless of other quality signals.
- Deduct 5 pts if webhook retries block the request thread instead of running in a background worker. Look for `HTTP:::retry` or equivalent being called from a controller.
- Deduct 5 pts if the condition tree is flattened (only supports a single AND at the top level, no nesting).
- Deduct 10 pts if any of the three TDD targets lacks a failing-test commit before the implementation commit. This is a hard rule; do not soften it because the student's final tests pass.
- If the retry test asserts only "was retried" without checking the interval schedule, treat the TDD category as Adequate rather than Excellent.
- Plugin evidence must show the plugin's output, not just a mention of the plugin's name. A screenshot of a slash command being invoked with the resulting output is the canonical form.